

OBJECTIVE:

preliminary prognosis of Hypertension/hypotension, based on the level of hemoglobin and genetic history of the individual.

```
In [2]: import pandas as pd
import numpy as np
df = pd.read_csv("data.csv")
```

```
In [3]: df.shape
```

Out[3]: (2000, 15)

```
In [4]: df.head()
```

Out[4]:

	Patient_Number	Blood_Pressure_Abnormality	Level_of_Hemoglobin	Genetic_Pedigree_Coefficient
0	1	1	11.28	0.90
1	2	0	9.75	0.23
2	3	1	10.79	0.91
3	4	0	11.00	0.43
4	5	1	14.17	0.83

```
In [5]: df.describe(include='all')
```

Out[5]:

	Patient_Number	Blood_Pressure_Abnormality	Level_of_Hemoglobin	Genetic_Pedigree_Coefficient
count	2000.000000	2000.000000	2000.000000	1908.000000
mean	1000.500000	0.493500	11.710035	0.493500
std	577.494589	0.500083	2.186701	0.293500
min	1.000000	0.000000	8.100000	0.000000
25%	500.750000	0.000000	10.147500	0.243500
50%	1000.500000	0.000000	11.330000	0.493500
75%	1500.250000	1.000000	12.945000	0.743500
max	2000.000000	1.000000	17.560000	1.000000

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 15 columns):
Patient_Number                2000 non-null int64
Blood_Pressure_Abnormality    2000 non-null int64
Level_of_Hemoglobin           2000 non-null float64
Genetic_Pedigree_Coefficient  1908 non-null float64
Age                           2000 non-null int64
BMI                            2000 non-null int64
Sex                            2000 non-null int64
Pregnancy                     442 non-null float64
Smoking                       2000 non-null int64
Physical_activity              2000 non-null int64
salt_content_in_the_diet      2000 non-null int64
alcohol_consumption_per_day   1758 non-null float64
Level_of_Stress                2000 non-null int64
Chronic_kidney_disease        2000 non-null int64
Adrenal_and_thyroid_disorders 2000 non-null int64
dtypes: float64(4), int64(11)
memory usage: 234.5 KB
```

Missing Value Treatment

```
In [7]: df.fillna(0)
```

Out[7]:

	Patient_Number	Blood_Pressure_Abnormality	Level_of_Hemoglobin	Genetic_Pedigree_Coeffici
0	1	1	11.28	C
1	2	0	9.75	C
2	3	1	10.79	C
3	4	0	11.00	C
4	5	1	14.17	C
5	6	0	11.64	C
6	7	1	11.69	C
7	8	0	12.70	C
8	9	0	10.88	C
9	10	1	14.56	C
10	11	1	8.58	C
11	12	1	12.77	C
12	13	1	16.40	C
13	14	0	16.42	C
14	15	0	11.97	C
15	16	1	10.96	C
16	17	1	11.98	C
17	18	1	11.60	C
18	19	1	8.99	C
19	20	1	16.55	C
20	21	1	16.95	C
21	22	0	10.85	C
22	23	0	11.19	C
23	24	0	10.61	C
24	25	0	10.26	C
25	26	1	13.29	C
26	27	0	10.40	C
27	28	0	10.55	C
28	29	1	9.37	C
29	30	0	11.17	C
...	
1970	1971	1	11.26	C
1971	1972	0	12.53	C
1972	1973	0	9.82	C

	Patient_Number	Blood_Pressure_Abnormality	Level_of_Hemoglobin	Genetic_Pedigree_Coeffici
1973	1974	1	11.45	C
1974	1975	1	8.47	C
1975	1976	1	8.75	C
1976	1977	0	12.07	C
1977	1978	0	11.16	C
1978	1979	0	11.57	C
1979	1980	0	11.70	C
1980	1981	0	12.34	C
1981	1982	1	13.49	C
1982	1983	0	9.08	C
1983	1984	1	14.56	C
1984	1985	0	10.81	C
1985	1986	0	10.44	C
1986	1987	1	13.63	C
1987	1988	1	11.79	C
1988	1989	1	17.17	C
1989	1990	1	10.84	C
1990	1991	1	11.21	C
1991	1992	1	15.53	C
1992	1993	1	9.38	C
1993	1994	0	9.69	1
1994	1995	0	11.07	C
1995	1996	1	10.14	C
1996	1997	1	11.77	1
1997	1998	1	16.91	C
1998	1999	0	11.15	C
1999	2000	1	11.36	C

2000 rows × 15 columns



```
In [8]: df.isnull().any(axis=0)
```

```
Out[8]: Patient_Number           False
Blood_Pressure_Abnormality      False
Level_of_Hemoglobin            False
Genetic_Pedigree_Coefficient    True
Age                             False
BMI                             False
Sex                             False
Pregnancy                      True
Smoking                        False
Physical_activity               False
salt_content_in_the_diet        False
alcohol_consumption_per_day     True
Level_of_Stress                 False
Chronic_kidney_disease          False
Adrenal_and_thyroid_disorders   False
dtype: bool
```

```
In [9]: df['Genetic_Pedigree_Coefficient'].isnull().value_counts()
```

```
Out[9]: False    1908
        True      92
        Name: Genetic_Pedigree_Coefficient, dtype: int64
```

```
In [10]: #median imputation
median = df['Genetic_Pedigree_Coefficient'].median()
df['Genetic_Pedigree_Coefficient'] = df['Genetic_Pedigree_Coefficient'].fillna(median)
```

```
In [11]: #checking for hidden missing values
print("# rows in dataframe {0}".format(len(df)))
print("# rows missing in Level_of_Hemoglobin: {0}".format(df['Level_of_Hemoglobin'].isnull().sum()))
print("# rows missing in Genetic_Pedigree_Coefficient: {0}".format(df['Genetic_Pedigree_Coefficient'].isnull().sum()))

# rows in dataframe 2000
# rows missing in Level_of_Hemoglobin: 0
# rows missing in Genetic_Pedigree_Coefficient: 0
```

```
In [12]: #check True/False ratio
num_false = len(df[df['Blood_Pressure_Abnormality'] == 0])
num_true = len(df[df['Blood_Pressure_Abnormality'] == 1])
print("no. of false cases: {0} ({1:2.2f}%)".format(num_false, (num_false*100/(num_false+num_true))))
print("no. of true cases: {0} ({1:2.2f}%)".format(num_true, (num_true*100/(num_false+num_true))))
#both cases have fair amount of cases standard predictions can be used
#Good distribution of true and false cases. No special work needed.


no. of false cases: 1013 (50.65%)
no. of true cases: 987 (49.35%)
```

```
In [13]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
sns.distplot(df['BMI'])
```

Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x23a4e4b8b70>


```
In [14]: sns.distplot(df['Physical_activity'])
```

Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x23a4fb7e860>



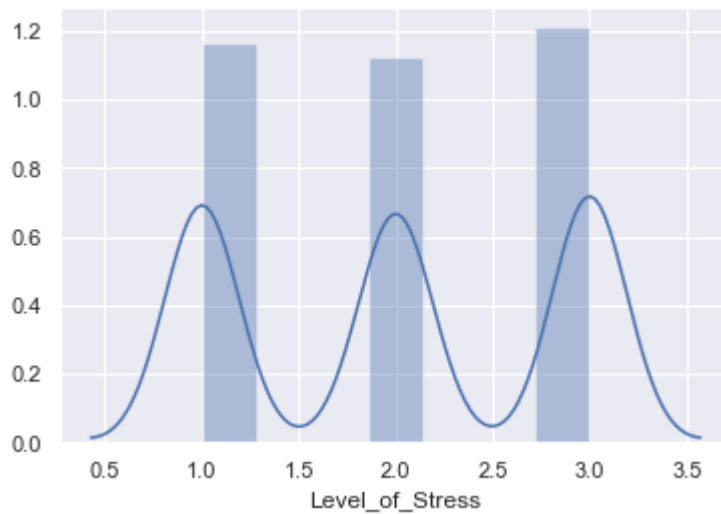
```
In [15]: sns.distplot(df['Smoking'])
```

Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x23a4fc299b0>



```
In [16]: sns.distplot(df['Level_of_Stress'])
```


```
Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x23a4ffcebe0>
```



```
In [17]: df['Genetic_Pedigree_Coefficient'].isnull().value_counts()
```

```
Out[17]: False      2000  
         Name: Genetic_Pedigree_Coefficient, dtype: int64
```

```
In [18]: sns.distplot(df['Genetic_Pedigree_Coefficient']);
```



```
In [19]: #print("Male Patient: {}".format(df[df["Sex"]== 0]['Sex'].count()))
# print("Female Patient: {}".format(df[df["Sex"]== 1]['Sex'].count()))
```

```
In [20]: #df[df["Sex"]== 1]['Pregnancy'].isnull().value_counts()
```

```
In [21]: #df[df["Sex"]== 0]['Pregnancy'].isnull().value_counts()
```

```
In [22]: df["Pregnancy"]=df["Pregnancy"].replace(df[df["Sex"]== 0]['Pregnancy'] ,0 )
```

```
In [23]: df["Pregnancy"]=df["Pregnancy"].replace(np.nan,0 )
df["Pregnancy"].isnull().value_counts()
```

```
Out[23]: False      2000
Name: Pregnancy, dtype: int64
```

Split Dataset(70,30)

```
In [25]: from sklearn.model_selection import train_test_split
feature_col_names = ['Genetic_Pedigree_Coefficient','BMI','Physical_activity','L
predicted_class_names = ['Blood_Pressure_Abnormality']
X= df[feature_col_names].values
y= df[predicted_class_names].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_
```

```
In [26]: #check to ensure we have the the desired 70% train, 30% test split of the data
print("{0:0.2f}% in training set".format((len(X_train)/len(df.index)) * 100))
print("{0:0.2f}% in test set".format((len(X_test)/len(df.index)) * 100))
```

```
70.00% in training set
30.00% in test set
```



```
In [27]: #Verifying predicted value was split correctly

print("Original True  : {0} ({1:0.2f}%)".format(len(df.loc[df['Blood_Pressure_Abnormality'] == 1]), (len(df.loc[df['Blood_Pressure_Abnormality'] == 1]) / len(df) * 100))
print("Original False : {0} ({1:0.2f}%)".format(len(df.loc[df['Blood_Pressure_Abnormality'] == 0]), (len(df.loc[df['Blood_Pressure_Abnormality'] == 0]) / len(df) * 100))
print("")
print("Training True   : {0} ({1:0.2f}%)".format(len(y_train[y_train[:] == 1]), (len(y_train[y_train[:] == 1]) / len(y_train) * 100))
print("Training False  : {0} ({1:0.2f}%)".format(len(y_train[y_train[:] == 0]), (len(y_train[y_train[:] == 0]) / len(y_train) * 100))
print("")
print("Test True       : {0} ({1:0.2f}%)".format(len(y_test[y_test[:] == 1]), (len(y_test[y_test[:] == 1]) / len(y_test) * 100))
print("Test False      : {0} ({1:0.2f}%)".format(len(y_test[y_test[:] == 0]), (len(y_test[y_test[:] == 0]) / len(y_test) * 100))
```

```
Original True   : 987 (49.35%)
Original False  : 1013 (50.65%)
```

```
Training True   : 707 (50.50%)
Training False  : 693 (49.50%)
```

```
Test True       : 280 (46.67%)
Test False      : 320 (53.33%)
```

NAIVE BAYES

```
In [28]: #Training the model
from sklearn.naive_bayes import GaussianNB

# create Gaussian Naive Bayes model object and train it with the data
nb_model = GaussianNB()

nb_model.fit(X_train, y_train.ravel())
```

```
Out[28]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [29]: # predict values using the training data
nb_predict_train = nb_model.predict(X_train)

# import the performance metrics library
from sklearn import metrics

# Accuracy
print("Accuracy on Training data: {0:.4f}".format(metrics.accuracy_score(y_train, nb_predict_train)))
print()
```

```
Accuracy on Training data: 0.5893
```

```
In [30]: # predict values using the testing data
nb_predict_test = nb_model.predict(X_test)

# Accuracy
print("Accuracy on Test data: {:.4f}".format(metrics.accuracy_score(y_test, nb_
print())
```

Accuracy on Test data: 0.5983

```
In [31]: score_nb1=metrics.accuracy_score(y_test,nb_predict_test)
print(score_nb1)
#print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
print("Accuracy on Test data: {:.4f}".format(metrics.accuracy_score(y_test, nb_
print())
```

0.5983333333333334

Accuracy on Test data: 0.5983

```
In [32]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,nb_predict_test)
```

```
Out[32]: array([[254,  66],
               [175, 105]], dtype=int64)
```

```
In [33]: from sklearn.metrics import classification_report
print(classification_report(y_test,nb_predict_test))
```

	precision	recall	f1-score	support
0	0.59	0.79	0.68	320
1	0.61	0.38	0.47	280
micro avg	0.60	0.60	0.60	600
macro avg	0.60	0.58	0.57	600
weighted avg	0.60	0.60	0.58	600

SUPPORT VECTOR MACHINE

```
In [34]: from sklearn.svm import SVC
model=SVC(gamma='auto')
```

```
In [35]: model.fit(X_train,y_train.ravel())
```

```
Out[35]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
In [36]: prediction_test=model.predict(X_test)
```

```
In [37]: print("accuracy on training data :{0:.4f}".format(model.score(X_test,y_test)))
```

accuracy on training data :0.4617

```
In [38]: prediction_test=model.predict(X_test)
score_svm1=metrics.accuracy_score(y_test,prediction_test)
print(score_svm1)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
```

0.46166666666666667

Accuracy= 0.46166666666666667

```
In [39]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[39]: array([[ 3, 317],
               [ 6, 274]], dtype=int64)
```

```
In [40]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.33	0.01	0.02	320
1	0.46	0.98	0.63	280
micro avg	0.46	0.46	0.46	600
macro avg	0.40	0.49	0.32	600
weighted avg	0.39	0.46	0.30	600

DECISION TREE

```
In [41]: from sklearn import tree
model=tree.DecisionTreeClassifier()
```

```
In [42]: model.fit(X_train,y_train.ravel())
```

```
Out[42]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
```

```
In [43]: prediction_test=model.predict(X_test)
```

```
In [44]: prediction_test=model.predict(X_test)
score_dt1=metrics.accuracy_score(y_test,prediction_test)
print(score_dt1)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
```

```
0.6466666666666666
Accuracy= 0.6466666666666666
```

```
In [45]: print("accuracy on training data :{0:.4f}".format(model.score(X_test,y_test)))
```

```
accuracy on training data :0.6467
```

```
In [46]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[46]: array([[212, 108],
               [104, 176]], dtype=int64)
```

```
In [47]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.67	0.66	0.67	320
1	0.62	0.63	0.62	280
micro avg	0.65	0.65	0.65	600
macro avg	0.65	0.65	0.65	600
weighted avg	0.65	0.65	0.65	600

RANDOM FOREST

```
In [48]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators = 10,random_state=30)
model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
from sklearn import metrics
score_rf1=(metrics.accuracy_score(y_test,prediction_test))
print(score_rf1)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
```

```
0.6933333333333334
Accuracy= 0.6933333333333334
```

```
In [49]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[49]: array([[243, 77],
               [107, 173]], dtype=int64)
```

```
In [50]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.69	0.76	0.73	320
1	0.69	0.62	0.65	280
micro avg	0.69	0.69	0.69	600
macro avg	0.69	0.69	0.69	600
weighted avg	0.69	0.69	0.69	600

GRADIENT BOOSTING

```
In [51]: from sklearn.ensemble import GradientBoostingClassifier
model=GradientBoostingClassifier(n_estimators = 10,random_state=30)
model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
from sklearn import metrics
score_gb1=format(model.score(X_test,y_test))
print("accuracy on testing data :{0:.4f}".format(score_gb1))
#print("accuracy on testing data :{0:.4f}".format(model.score(X_test,y_test)))
print("accuracy on testing data :{0:.4f}".format(model.score(X_test,y_test)))
```

```
accuracy on testing data :{0:.4f} 0.7566666666666667
accuracy on testing data :0.7567
```

```
In [52]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[52]: array([[292,  28],
               [118, 162]], dtype=int64)
```

```
In [53]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.71	0.91	0.80	320
1	0.85	0.58	0.69	280
micro avg	0.76	0.76	0.76	600
macro avg	0.78	0.75	0.74	600
weighted avg	0.78	0.76	0.75	600

```
In [54]: scores = [score_nb,score_svm,score_dt,score_rf,score_gb]
          algorithms = ["Naive Bayes","Support Vector Machine","Decision Tree","Random Forest"]

          for i in range(len(algorithms)):
              print("The accuracy score achieved using "+algorithms[i]+" is: "+str(scores[i]))
```

The accuracy score achieved using Naive Bayes is: 0.5983333333333334 %

The accuracy score achieved using Support Vector Machine is: 0.4616666666666667 %

The accuracy score achieved using Decision Tree is: 0.6466666666666666 %

The accuracy score achieved using Random Forest is: 0.6933333333333334 %


The accuracy score achieved using XGBoost is: 0.7566666666666667 %

```
In [98]: scores = [score_nb,score_svm,score_dt,score_rf,score_gb]
          algorithms = ["Naive Bayes","Support Vector Machine","Decision Tree","Random Forest"]
```

```
In [99]: sns.set(rc={'figure.figsize':(15,8)})
          plt.xlabel("Algorithms")
          plt.ylabel("Accuracy score")

          sns.lineplot(algorithms,scores)
```


Out[99]: <matplotlib.axes._subplots.AxesSubplot at 0x23a50131278>



```
In [100]: sns.set(rc={'figure.figsize':(15,8)})
plt.xlabel("Algorithms")
plt.ylabel("Accuracy score")

sns.barplot(algorithms,scores)
```

```
Out[100]: <matplotlib.axes._subplots.AxesSubplot at 0x23a5085e240>
```



Split Dataset

```
In [55]: from sklearn.model_selection import train_test_split
feature_col_names = ['Level_of_Hemoglobin', 'Genetic_Pedigree_Coefficient']
predicted_class_names = ['Blood_Pressure_Abnormality']
X= df[feature_col_names].values
y= df[predicted_class_names].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random
```

NAIVE BAYES

```
In [56]: #Training the model
from sklearn.naive_bayes import GaussianNB

# create Gaussian Naive Bayes model object and train it with the data
nb_model = GaussianNB()

nb_model.fit(X_train, y_train.ravel())
```

Out[56]: GaussianNB(priors=None, var_smoothing=1e-09)

```
In [57]: # predict values using the testing data
nb_predict_test = nb_model.predict(X_test)

# Accuracy
print("Accuracy on Test data: {0:.4f}".format(metrics.accuracy_score(y_test, nb_
print()
```

Accuracy on Test data: 0.8657

```
In [58]: nb_predict_test = nb_model.predict(X_test)

# Accuracy
score_nb=metrics.accuracy_score(y_test,nb_predict_test)
print(score_nb)
#print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
print("Accuracy on Test data: {0:.4f}".format(metrics.accuracy_score(y_test, nb_
print()
```

0.8657142857142858

Accuracy on Test data: 0.8657

```
In [59]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,nb_predict_test)
```

Out[59]: array([[322, 45],
[49, 284]], dtype=int64)

```
In [60]: from sklearn.metrics import classification_report
print(classification_report(y_test,nb_predict_test))
```

	precision	recall	f1-score	support
0	0.87	0.88	0.87	367
1	0.86	0.85	0.86	333
micro avg	0.87	0.87	0.87	700
macro avg	0.87	0.87	0.87	700
weighted avg	0.87	0.87	0.87	700

SUPPORT VECTOR MACHINE


```
In [61]: from sklearn.svm import SVC
model=SVC(gamma='auto')
```

```
In [62]: model.fit(X_train,y_train.ravel())
```

```
Out[62]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [63]: prediction_test=model.predict(X_test)
score_svm=metrics.accuracy_score(y_test,prediction_test)
print(score_svm)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
```

```
0.8628571428571429
Accuracy= 0.8628571428571429
```

```
In [64]: print("accuracy on testing data :{0:.4f}".format(model.score(X_test,y_test)))
```

```
accuracy on testing data :0.8629
```

```
In [65]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[65]: array([[335,  32],
[ 64, 269]], dtype=int64)
```

```
In [66]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.84	0.91	0.87	367
1	0.89	0.81	0.85	333
micro avg	0.86	0.86	0.86	700
macro avg	0.87	0.86	0.86	700
weighted avg	0.87	0.86	0.86	700

DECISION TREE

```
In [67]: from sklearn import tree
model=tree.DecisionTreeClassifier()
```

```
In [68]: model.fit(X_train,y_train.ravel())
```

```
Out[68]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [69]: prediction_test=model.predict(X_test)
         score_dt=metrics.accuracy_score(y_test,prediction_test)
         print(score_dt)
         print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
```

```
0.8514285714285714
Accuracy= 0.8514285714285714
```

```
In [70]: print("accuracy on testing data :{0:.4f}".format(model.score(X_test,y_test)))
```

```
accuracy on testing data :0.8514
```

```
In [71]: from sklearn.metrics import confusion_matrix
         confusion_matrix(y_test,prediction_test)
```

```
Out[71]: array([[322,  45],
                [ 59, 274]], dtype=int64)
```

```
In [72]: from sklearn.metrics import classification_report
         print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.85	0.88	0.86	367
1	0.86	0.82	0.84	333
micro avg	0.85	0.85	0.85	700
macro avg	0.85	0.85	0.85	700
weighted avg	0.85	0.85	0.85	700

RANDOM FOREST

```
In [73]: from sklearn.ensemble import RandomForestClassifier
         model = RandomForestClassifier(n_estimators = 10,random_state=30)
         model.fit(X_train,y_train.ravel())
         prediction_test=model.predict(X_test)
         from sklearn import metrics
         scores_rf=(metrics.accuracy_score(y_test,prediction_test))
         print("Accuracy on test data=",scores_rf)
```

```
Accuracy on test data= 0.9057142857142857
```

```
In [74]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[74]: array([[342, 25],
               [ 41, 292]], dtype=int64)
```

```
In [75]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

	precision	recall	f1-score	support
0	0.89	0.93	0.91	367
1	0.92	0.88	0.90	333
micro avg	0.91	0.91	0.91	700
macro avg	0.91	0.90	0.91	700
weighted avg	0.91	0.91	0.91	700

GRADIENT BOOSTING

```
In [76]: from sklearn.ensemble import GradientBoostingClassifier
model=GradientBoostingClassifier(n_estimators = 10,random_state=30)
model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
from sklearn import metrics
score_gb=format(model.score(X_test,y_test))
print("accuracy on testing data :{0:.4f}",score_gb)
print("accuracy on testing data :{0:.4f}".format(model.score(X_test,y_test)))
```

```
accuracy on testing data :{0:.4f} 0.9
accuracy on testing data :0.9000
```

```
In [77]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[77]: array([[342, 25],
               [ 45, 288]], dtype=int64)
```

```
In [78]: from sklearn.metrics import classification_report
print(classification_report(y_test,prediction_test))
```

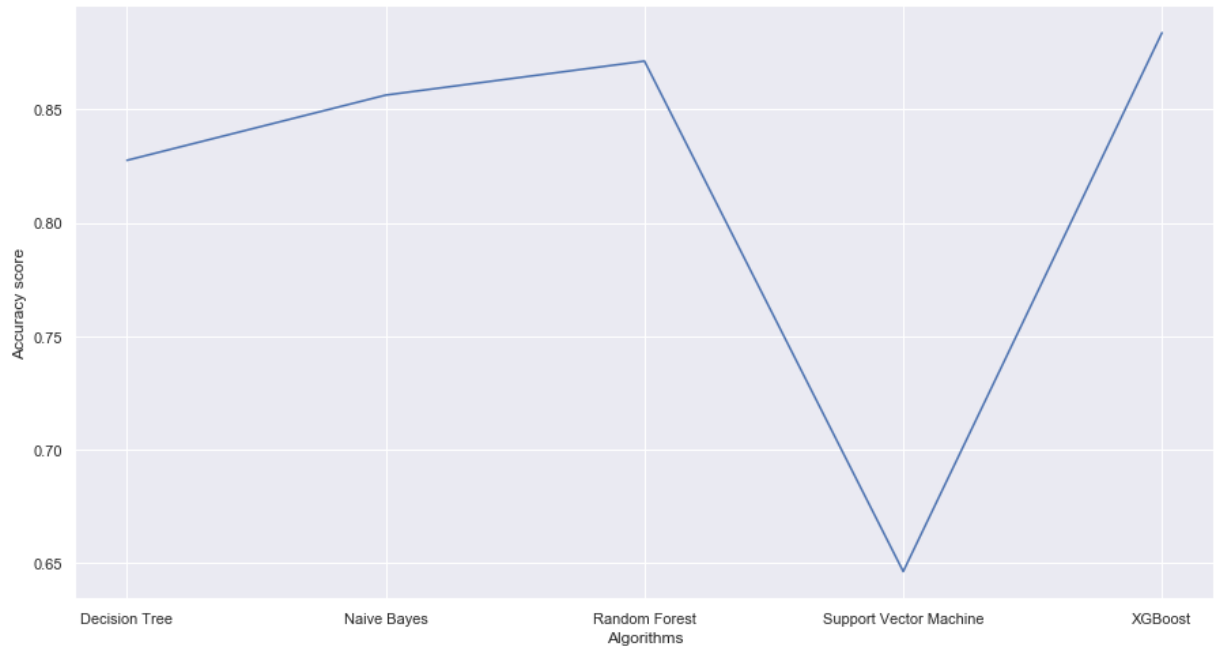
	precision	recall	f1-score	support
0	0.88	0.93	0.91	367
1	0.92	0.86	0.89	333
micro avg	0.90	0.90	0.90	700
macro avg	0.90	0.90	0.90	700
weighted avg	0.90	0.90	0.90	700

```
In [101]: scores = [score_nb,score_svm,score_dt,score_rf,score_gb]
          algorithms = ["Naive Bayes","Support Vector Machine","Decision Tree","Random Forest"]
```

```
In [102]: sns.set(rc={'figure.figsize':(15,8)})
          plt.xlabel("Algorithms")
          plt.ylabel("Accuracy score")

          sns.lineplot(algorithms,scores)
```


```
Out[102]: <matplotlib.axes._subplots.AxesSubplot at 0x23a51900ac8>
```



```
In [103]: sns.set(rc={'figure.figsize':(15,8)})
plt.xlabel("Algorithms")
plt.ylabel("Accuracy score")

sns.barplot(algorithms,scores)
```

```
Out[103]: <matplotlib.axes._subplots.AxesSubplot at 0x23a52213f60>
```



splitdata

```
In [79]: from sklearn.model_selection import train_test_split
feature_col_names = ['Level_of_Hemoglobin', 'Genetic_Pedigree_Coefficient', 'Age']
predicted_class_names = ['Blood_Pressure_Abnormality']
X= df[feature_col_names].values
y= df[predicted_class_names].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_
```

NAIVE BAYES

```
In [80]: #Training the model
from sklearn.naive_bayes import GaussianNB

# create Gaussian Naive Bayes model object and train it with the data
nb_model = GaussianNB()

nb_model.fit(X_train, y_train.ravel())
```

```
Out[80]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [81]: # predict values using the testing data
nb_predict_test = nb_model.predict(X_test)

# Accuracy
print("Accuracy on Test data: {0:.4f}".format(metrics.accuracy_score(y_test, nb_
print()
```

Accuracy on Test data: 0.8562

```
In [82]: nb_predict_test = nb_model.predict(X_test)

# Accuracy
score_nb=metrics.accuracy_score(y_test,nb_predict_test)
print(score_nb)
#print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
print("Accuracy on Test data: {0:.4f}".format(metrics.accuracy_score(y_test, nb_
print()
```

0.85625

Accuracy on Test data: 0.8562

```
In [83]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,nb_predict_test)
```

```
Out[83]: array([[368,  58],
               [ 57, 317]], dtype=int64)
```

```
In [84]: from sklearn.metrics import classification_report
print(classification_report(y_test,nb_predict_test))
```

	precision	recall	f1-score	support
0	0.87	0.86	0.86	426
1	0.85	0.85	0.85	374
micro avg	0.86	0.86	0.86	800
macro avg	0.86	0.86	0.86	800
weighted avg	0.86	0.86	0.86	800

SUPPORT VECTOR MACHINE

```
In [85]: from sklearn.svm import SVC
model=SVC(gamma='auto')
model.fit(X_train,y_train.ravel())
```

```
Out[85]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

```
In [86]: prediction_test=model.predict(X_test)
```

```
In [87]: score_svm=metrics.accuracy_score(y_test,prediction_test)
print(score_svm)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))

0.64625
Accuracy= 0.64625
```

```
In [88]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,prediction_test)
```

```
Out[88]: array([[267, 159],
               [124, 250]], dtype=int64)
```

DECISION TREE

```
In [89]: from sklearn import tree
model=tree.DecisionTreeClassifier()
```

```
In [90]: model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
```

```
In [91]: score_dt=metrics.accuracy_score(y_test,prediction_test)
print(score_dt)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))

0.8275
Accuracy= 0.8275
```

```
In [92]: print("accuracy on training data :{0:.4f}".format(model.score(X_test,y_test)))

accuracy on training data :0.8275
```

RANDOM FOREST

```
In [93]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier(n_estimators = 10,random_state=30)
model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
from sklearn import metrics
score_rf=metrics.accuracy_score(y_test,prediction_test)
print(score_rf)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))

0.87125
Accuracy= 0.87125
```

GRADIENT BOOSTING

```
In [94]: from sklearn.ensemble import GradientBoostingClassifier
model=GradientBoostingClassifier(n_estimators = 10,random_state=30)
model.fit(X_train,y_train.ravel())
prediction_test=model.predict(X_test)
from sklearn import metrics
score_gb=metrics.accuracy_score(y_test,prediction_test)
print(score_gb)
print("Accuracy=",metrics.accuracy_score(y_test,prediction_test))
print("accuracy on training data :{0:.4f}".format(model.score(X_test,y_test)))
```

0.88375

Accuracy= 0.88375


accuracy on training data :0.8838

```
In [95]: scores = [score_nb,score_svm,score_dt,score_rf,score_gb]
algorithms = ["Naive Bayes","Support Vector Machine","Decision Tree","Random Forest"]
```

```
In [96]: sns.set(rc={'figure.figsize':(15,8)})
plt.xlabel("Algorithms")
plt.ylabel("Accuracy score")

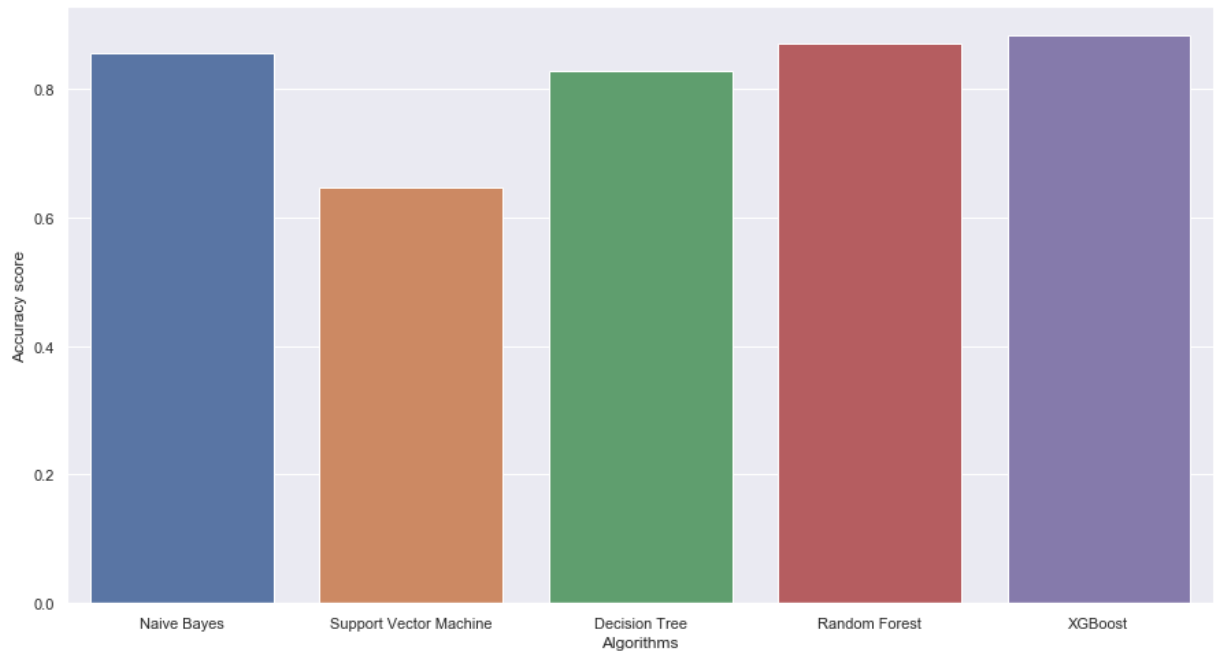
sns.lineplot(algorithms,scores)
```

Out[96]: <matplotlib.axes._subplots.AxesSubplot at 0x23a519d67f0>




```
In [97]: sns.set(rc={'figure.figsize':(15,8)})  
plt.xlabel("Algorithms")  
plt.ylabel("Accuracy score")  
  
sns.barplot(algorithms,scores)
```

Out[97]: <matplotlib.axes._subplots.AxesSubplot at 0x23a501a3dd8>



In []: